

RACHEL FREEDMAN

rachel.freedman@berkeley.edu | github.com/RachelFreedman | <https://rachelfreedman.github.io/>

Education

2019-2027 (expected)	Artificial Intelligence PhD, UC Berkeley I work with the Center for Human-Compatible Artificial Intelligence (CHAI) on AI safety via reinforcement learning, reward modeling, and interpretability. I am a Cooperative AI Foundation PhD Fellow and my advisor is Professor Stuart Russell.	3.93 GPA
2013-2017	Computer Science/Psychology BA, Duke University • Graduated Phi Beta Kappa and <i>magna cum laude</i> • Earned High Distinction in Computer Science for my research thesis (published in the journal <i>Artificial Intelligence</i>) • Designed original interdepartmental major on "AI Systems"	3.93 GPA
2015-2016	CS/Philosophy Registered Visiting Student, Oxford University Founded AI safety student society at Future of Humanity Institute	3.93 GPA
2011-2017	Part-Time Student, UNC Chapel Hill	4.00 GPA

Publications

Rachel Freedman, Justin Svegliato, Kyle Wray, Stuart Russell. "Active Teacher Selection for Reward Learning". In *TMLR*, contributed talk at *PRL at AAAI 2025* and *TAIS 2025*.

Rachel Freedman. "Adaptive Pluralistic Alignment: A Pipeline for Dynamic Artificial Democracy". Under review at *Pluralistic Alignment Workshop at ICML 2026*.

Henry Papadatos, **Rachel Freedman**. "Linear Probe Penalties Reduce LLM Sycophancy". In *Socially Responsible Language Modelling Research Workshop (SoLaR) at NeurIPS 2024*.

Vincent Conitzer, **Rachel Freedman**, Jobst Heitzig, Wesley H. Holliday, Bob M. Jacobs, Nathan Lambert, Milan Mossé, Eric Pacuit, Stuart Russell, Hailey Schoelkopf, Emanuel Tewolde, William S. Zwicker. "Social Choice for AI Alignment: Dealing with Diverse Human Feedback". In *ICML 2024*.

Peter Barnett, **Rachel Freedman**, Justin Svegliato, Stuart Russell. "Active Reward Learning from Multiple Teachers". In *SafeAI at AAAI 2023*.

Best Paper Award Finalist at AAAI 2023 SafeAI Workshop.

Stephen Casper, [...], **Rachel Freedman**, [...], Dylan Hadfield-Menell. "Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback". In *TMLR*.

Oliver Daniels-Koch, **Rachel Freedman**. "The Expertise Problem: Learning from Specialized Feedback". In *ML Safety Workshop at NeurIPS 2022*.

AI Risk Analysis Award at NeurIPS 2022 ML Safety Workshop.

Rachel Freedman, Jana Schaich Borg, Walter Sinnott-Armstrong, John Dickerson, and Vincent Conitzer. "Adapting a Kidney Exchange Algorithm to Align with Human Values". *Artificial Intelligence*, v. 283, 2020. Also presented at *AAAI 2018*, *AIES 2018*, *MD4SG 2018*, and the *Participatory ML workshop at ICML 2020*.

Outstanding Student Paper Honorable Mention at AAAI 2018.

Rachel Freedman, Rohin Shah, and Anca Dragan. "Choice Set Misspecification in Reward Inference". In *Workshop on Artificial Intelligence Safety at IJCAI 2020*.

Best Paper Award at IJCAI 2020 AISafety Workshop.

Rohin Shah, Pedro Freire, Neel Alex, **Rachel Freedman**, Dmitrii Krasheninnikov, Lawrence Chan, Michael Dennis, Pieter Abbeel, Anca Dragan and Stuart Russell. "Benefits of Assistance over Reward Learning". In *Cooperative AI Workshop at NeurIPS 2020*.

Best Paper Award at NeurIPS 2020 CoopAI Workshop.

Work and Research

2019-pres.	Berkeley PhD Researcher <i>CHAI lab (UC Berkeley)</i> • Research topics include reinforcement learning, reward modeling, interpretability, LLM capability evaluation, democratic AI, and pluralistic alignment in support of robustly beneficial and safe AI • Contributed technical feedback to reports for the United Nations and the World Economic Forum and to books including <i>The Alignment Problem</i> (by Brian Christian) and <i>Moral AI</i> (by Borg, Conitzer and Sinnott-Armstrong) • Advised by Prof. Stuart Russell, also worked with Prof. Anca Dragan
2023	Cambridge Visiting Researcher <i>Krueger lab (Cambridge University)</i> • Collaborated with Prof. David Krueger, Computational and Biological Learning Lab • Researched reducing causal confusion in reward learning from human feedback
2017-2019	Software Engineer and Consultant <i>Galatea Associates (London, UK)</i> • Collaborated with international team of consultants, engineers, contractors and stakeholders to manage ever-changing business problems and constraints • Learned new tools and languages rapidly as required
2016-2017	Moral AI Researcher <i>Moral AI lab (Duke University)</i> • Conducted a year-long independent research project incorporating computer science and psychology methodologies into the field of artificial morality • Adapted a kidney exchange algorithm to align with human values by modeling underlying frameworks in human responses to moral dilemmas • Awarded High Distinction for my thesis research; presented and published this work at AAAI 2018, where it won Outstanding Student Paper Honorable Mention • Published in <i>Artificial Intelligence</i>, poster presented at AIES 2018 and MD4SG 2018.
2015 summer	Software Engineering Intern <i>Microsoft (Seattle, US)</i>
2013-2014	Psychology and Neuroscience Research Assistant <i>Duke University</i>

Honors

Conferences & Workshops	Best Paper Award Finalist <i>SafeAI Workshop at AAAI (2023)</i> AI Risk Analysis Award <i>ML Safety Workshop at NeurIPS (2022)</i> Ambassador <i>Effective Altruism Global (2020 - 2021)</i> Best Paper Award <i>CoopAI Workshop at NeurIPS (2020)</i> Best Paper Award <i>AI Safety Workshop at IJCAI-PRICAI (2020)</i> Outstanding Student Paper Honorable Mention <i>AAAI (2018)</i>
Graduate	Cooperative AI PhD Fellowship <i>Cooperative AI Foundation (2025)</i> Long Term Future Fund Grant <i>EA Funds (2024)</i> Foresight Fellowship <i>Foresight Institute (2024)</i> Manifund Grant <i>Manifund (2023)</i> Rising Star in AI Ethics <i>Women in AI Ethics (2021)</i> EECS Excellence Award <i>UC Berkeley (2019)</i> EECS Departmental Fellowship <i>UC Berkeley (2019)</i>
Undergraduate	Phi Beta Kappa <i>Duke University (2017)</i> magna cum laude <i>Duke University (2017)</i> High Distinction in Computer Science <i>Duke University (2017)</i> Robertson Scholarship <i>Robertson Scholarship Leadership Program (2013-2017)</i> Thomas J. Watson Memorial Scholarship <i>IBM (2013-2016)</i>
Other	Rhodes Scholarship Finalist <i>Rhodes Trust (2017)</i> Gates Cambridge Finalist <i>Gates Cambridge Trust (2019)</i> President's Volunteer Service Gold Award <i>US Government (2013)</i>

Service and Leadership

2026	Program Committee <i>AAAI 2026 AI Alignment Track (AAAI-AIA26)</i>
2025	Conference Reviewer <i>International Association for Safe and Ethical AI (IASEAI26)</i>
2025	Program Committee <i>AAAI 2025 AI Alignment Track (AAAI-AIA25)</i>
2024	Conference Reviewer <i>Neural Information Processing Systems (NeurIPS)</i>
2024	Program Committee <i>ECAI Workshop on Social Choice for AI Ethics and Safety</i>
2022-pres	AI Safety Advisor <i>80000 Hours (contractor)</i>
2022	Workshop Reviewer <i>NeurIPS ML Safety Workshop (MLSW22)</i>
2022	Journal Reviewer <i>Journal of Artificial Intelligence (JAIR)</i>
2021	Program Committee <i>IJCAI Workshop on Artificial Intelligence Safety (NYC, US)</i>
2019-2021	Ambassador <i>Effective Altruism Global Conferences (London, UK and San Francisco, US)</i>
2015-2016	Co-founder and President <i>Oxford Existential Risk Society (Oxford, UK)</i>

Invited Talks and Panels

2026	Introduction to AI Safety and Ethics <i>Social Choice for AI Workshop</i>
2026	Reward Model Personalization <i>Center for Human-Compatible AI</i>
2025	Misspecification Failures in RLHF <i>London Initiative for Safe AI</i>
2025	Social Choice for AI Alignment <i>Center for Human-Compatible AI</i>
2024	Limitations of RLHF <i>Stanford CS362 guest lecture</i>
2024	Aligning AI to Diverse Principals <i>Foresight Institute Vision Weekend</i>
2024	Social Choice for AI Alignment <i>Institutional Economics (SIOE) Conference</i>
2024	Social Choice for RLHF <i>Cooperative AI Summer School lecture</i>
2024	Reinforcement Learning from Human Feedback <i>Stanford E338 guest lecture</i>
2024	RLHF: Problems and Approaches <i>Stanford Existential Risk Symposium</i>
2024	Open Problems in RLHF <i>ATHENA seminar series</i>
2023	Panel on AI and Business <i>Growth Factory Venture Conference</i>
2023	Active Teacher Selection <i>Cambridge University Krueger Lab</i>
2022	Value Alignment <i>Institute for Advanced Study WAM program</i>
2022	My Approach to Alignment Research <i>SERI MATS seminar series</i>
2022	Reward Modeling and Human Model Misspecification <i>CHAI Workshop</i>
2022	Panel on AI <i>Duke University Career Center</i>
2021	Panel on Graduate School <i>AI Safety Support</i>

Advising and Mentoring

2024-2025	Pavel Czempin	CHAI Intern	RSS 2025 Workshop CRLH
2023-2024	Henry Papadatos	EPFL MSc	NeurIPS SoLaR Workshop
2022-2023	Peter Barnett	CHAI Intern	AAAI Workshop (Award Finalist)
2022-2023	Oliver Daniels-Koch	CHAI Intern	NeurIPS Workshop (Award)

Skills and Hobbies

Programming	Python, NumPy, PyTorch, Java, Julia, Git; extensive practice learning new technologies quickly and proactively via engineering and consulting work
Service	Mentor young people seeking a career in AI (particularly those from underrepresented backgrounds) on an ongoing basis (4+ years), volunteer at the Berkeley public animal shelter (2 years), tutored and mentored at-risk North Carolina students weekly (2 years, earned the President's Volunteer Service Gold Award), volunteered full-time to support students in rural Mississippi (3 months).
Martial Arts	Earned World Taekwondo Federation-certified master black belt (4th degree) in Taekwondo and 2nd degree in Kumdo. Trained, taught and competed for 12 years.
Creative Writing	Served as writer-in-residence at artist residency <i>Arte Studio Ginestrelle</i> in Italy. Stories and plays published in Duke, UNC and Oxford University literary publications, and performed at Duke New Works Festival.

Last Updated April 2026