RACHEL FREEDMAN

rachel.freedman@berkeley.edu | github.com/RachelFreedman | https://rachelfreedman.github.io/

Education

late 2019 - present *(expected May 2027)*	**Artificial Intelligence PhD, *UC Berkeley*** I work with the Center for Human-Compatible Artificial Intelligence (CHAI) on AI safety via reinforcement learning, reward modeling, and interpretability. I am a **Cooperative AI Foundation PhD Fellow** advised by Professor Stuart Russell.	**3.93 GPA**
2013-2017	Computer Science/Psychology BA, ***Duke University*** • Graduated **Phi Beta Kappa** and ***magna cum laude*** • Earned **High Distinction** in Computer Science for my research thesis (published in the journal *Artificial Intelligence*)	**3.93 GPA**
2015-2016	CS/Philosophy Visiting Student, ***Oxford University*** I founded the AI safety student society at Future of Humanity Institute.	**3.93 GPA**

Publications

Rachel Freedman, Justin Svegliato, Kyle Wray, Stuart Russell. "Active Teacher Selection for Reward Learning". In *TMLR*, contributed talk at *PRL at AAAI 2025* and *TAIS 2025*.

Rachel Freedman. "Adaptive Pluralistic Alignment: A Pipeline for Dynamic Artificial Democracy". In *Pluralistic Alignment Workshop* at *ICML 2026.*

Henry Papadatos, **Rachel Freedman**. "Linear Probe Penalties Reduce LLM Sycophancy". In *Socially Responsible Language Modelling Research Workshop (SoLaR) at NeurIPS 2024.*

Vincent Conitzer, **Rachel Freedman**, Jobst Heitzig, Wesley H. Holliday, Bob M. Jacobs, Nathan Lambert, Milan Mossé, Eric Pacuit, Stuart Russell, Hailey Schoelkopf, Emanuel Tewolde, William S. Zwicker. "Social Choice for AI Alignment: Dealing with Diverse Human Feedback". In *ICML 2024.*

Peter Barnett, **Rachel Freedman**, Justin Svegliato, Stuart Russell. "Active Reward Learning from Multiple Teachers". In *SafeAI at AAAI 2023*. **✱ Best Paper Award Finalist ✱**

Stephen Casper, [...], **Rachel Freedman**, [...], Dylan Hadfield-Menell. "Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback". In *ICLR 2025* and *TMLR*.

Oliver Daniels-Koch, **Rachel Freedman**. "The Expertise Problem: Learning from Specialized Feedback". In *ML Safety Workshop at NeurIPS* 2022. **✱ AI Risk Analysis Award ✱**

Rachel Freedman, Jana Schaich Borg, Walter Sinnott-Armstrong, John Dickerson, and Vincent Conitzer. "Adapting a Kidney Exchange Algorithm to Align with Human Values". *Artificial Intelligence*, v. 283, 2020. Also presented at *AAAI 2018, AIES 2018*, *MD4SG 2018*, and the *Participatory ML workshop* at *ICML 2020*. **✱ Outstanding Student Paper Honorable Mention ✱** at *AAAI 2018*.

Rachel Freedman, Rohin Shah, and Anca Dragan. "Choice Set Misspecification in Reward Inference". In *Workshop on Artificial Intelligence Safety at IJCAI* 2020. **✱ Best Paper Award ✱**

Rohin Shah, Pedro Freire, Neel Alex, **Rachel Freedman**, Dmitrii Krasheninnikov, Lawrence Chan, Michael Dennis, Pieter Abbeel, Anca Dragan and Stuart Russell. "Benefits of Assistance over Reward Learning". In *Cooperative AI Workshop at NeurIPS* 2020. **✱ Best Paper Award ✱**

Work and Research

2019-pres. **Berkeley PhD Researcher** *CHAI lab (UC Berkeley)*
- Research topics include reward modeling (RLHF), interpretability, LLM capability evaluation, democratic AI, and pluralistic alignment in support of robustly beneficial and safe AI
- Mentored 4 interns, who have gone to lead AI safety non-profits (Henry Papadatos, SaferAI), work at AI safety non-profits (Peter Barnett, MIRI), and pursue PhDs in the

area (Pavel Czempin, USC; Oliver Daniels, UMass Amherst)
- Contributed technical feedback to reports for the United Nations and the World Economic Forum and to books including *The Alignment Problem* (by Brian Christian) and *Moral AI* (by Borg, Conitzer and Sinnott-Armstrong)
- Guest lectured multiple Stanford AI courses

2022-pres. **AI Safety Career Advisor** (contractor) *80,000 Hours (Remote)*
- Advised >200 participants across 5 continents on AI safety research and careers
- Conducted regular group calls to answer questions from global participants

2023 **Cambridge Visiting Researcher** *Prof. David Krueger's lab (Cambridge University)*
- Researched reducing causal confusion in reward learning from human feedback

2017-2019 **Software Engineer and Consultant** *Galatea Associates (London, UK)*
- Collaborated with international team of consultants, engineers, contractors and stakeholders to manage volatile business problems and constraints

2016-2017 **Moral AI Researcher** *Moral AI lab (Duke University)*
- Conducted a year-long independent research project incorporating computer science and psychology methodologies into the field of artificial morality
- Awarded High Distinction for my thesis research; presented and published this work at AAAI 2018, where it won Outstanding Student Paper Honorable Mention
- Published in *Artificial Intelligence*, poster presented at AIES 2018 and MD4SG 2018.

2015 summer **Software Engineering Intern** *Microsoft (Seattle, US)*

Honors

Conferences & Workshops
Best Paper Award Finalist *SafeAI Workshop at AAAI (2023)*
AI Risk Analysis Award *ML Safety Workshop at NeurIPS (2022)*
Best Paper Award *CoopAI Workshop at NeurIPS (2020)*
Best Paper Award *AISafety Workshop at IJCAI-PRICAI (2020)*
Outstanding Student Paper Honorable Mention *AAAI (2018)*

Graduate
Cooperative AI PhD Fellowship *Cooperative AI Foundation (2025)*
Long Term Future Fund Grant *EA Funds (2024)*
Foresight Fellowship *Foresight Institute (2024)*
Rising Star in AI Ethics *Women in AI Ethics (2021)*
EECS Departmental Fellowship *UC Berkeley (2019)*

Undergraduate
Phi Beta Kappa *Duke University (2017)*
magna cum laude *Duke University (2017)*
High Distinction in Computer Science *Duke University (2017)*
Robertson Scholarship *Robertson Scholarship Leadership Program (2013-2017)*

Service and Leadership

2026 **Grant Advisor** *Schmidt Sciences*

2026 **Stream Mentor** *PRISM AI Safety Fellowship*

2025-6 **Program Committee** *AAAI 2025 and 2026 AI Alignment Track*

2022-pres **Reviewer** *JAIR, NeurIPS, IASEAI, NeurIPS ML Safety Workshop*

2022-pres **AI Safety Advisor** *80000 Hours (contractor)*

2015-2016 **Co-founder and President** *Oxford Existential Risk Society (Oxford, UK)*